



Fusion Platform[®]

Python SDK

Issue 1 Draft
8th March 2022

Table of Contents

1	Introduction	4
1.1	Purpose.....	4
1.2	Scope.....	4
1.3	Overview	4
2	Overview	5
2.1	User Accounts, Organisations and Credits	5
2.2	Workflow	5
2.3	Installation	7
2.4	Usage.....	7
3	Quick Example.....	8
4	Authentication	10
5	Uploading Files	12
6	Browsing Services	16
7	Creating Processes.....	18
8	Executing and Monitoring Processes	21
9	Obtaining Execution Inputs and Outputs	25
10	Detailed Example.....	27
11	SDK Object Attributes and Methods.....	30
11.1	Attributes and Methods.....	30
12	Glossary of Terms	32
13	References	34

Revision History

Draft	7 th February 2022	Initial draft for review.
Draft	9 th February 2022	Updated following internal review.
Draft	22 nd February 2022	Updated during implementation.
Draft	3 rd March 2022	Issue candidate.
Draft	8 th March 2022	Revised following internal review to include link to code documentation.

This document has been prepared by Digital Content Analysis Technology Ltd, who can be contacted at support@d-cat.co.uk

Digital Content Analysis Technology Ltd, Registered in Scotland, SC499652.

This document is copyright and is issued on the strict understanding that it is not to be reproduced, copied, or disclosed to a third party, in whole or in part, without the consent of Digital Content Analysis Technology Ltd.

1 Introduction

1.1 Purpose

This document describes the Python SDK used to interact with the Fusion Platform®. The Fusion Platform® provides enhanced remote monitoring services. By ingesting remotely sensed Earth Observation (EO) data, and data from other sources, the platform uses and fuses this data to execute algorithms which provide actionable knowledge to customers.

1.2 Scope

The Python SDK is designed to enable interaction with the Fusion Platform® via its API. As such, the SDK therefore allows software to login, upload files, create and execute processes, monitor their execution and then download the corresponding results. Additional functionality is available directly via the API, and this is defined within the corresponding OpenAPI 3.0 specification, which can be obtained via a support request.

Support for the SDK can be obtained by contacting Digital Content Analysis Technology Ltd via support@d-cat.co.uk.

1.3 Overview

This document first presents a brief overview of the Fusion Platform® and the general requirements for the SDK. This is followed by an example of how to get started with the SDK for those who wish to try things out. Subsequent sections describe the functionality in detail. Finally, a more detailed example is provided, along with a full list of the SDK methods available, together with a glossary of terms.

Python code snippets within this document are shown as follows. These assume a basic level of Python knowledge. Each block is a snippet of code which may rely upon previous code examples to be completed to successfully execute.

Python

```
# This is a Python code snippet.  
print('Hello world!')
```

Where appropriate, shell commands are also shown. These should work regardless of the host operating system (for example, Windows PowerShell, Mac Terminal or Linux command line):

Command

```
echo "Hello world!"
```

When there are important notes or warnings, they are shown like this:

Important

This is an important note!

2 Overview

The Fusion Platform® is a cloud-based service which allows users to execute algorithms published by algorithm providers. Execution may be on-demand or at regular intervals using a variety of available EO and other data inputs. This data is automatically ingested by the platform to satisfy execution requirements. Users pay for execution and storage of files through pre-purchased credits.

2.1 User Accounts, Organisations and Credits

All access to the Fusion Platform® is via a registered user account. Account registration is performed separately to the SDK and must therefore be completed before the SDK can be used.

All user accounts are linked to an email address which can receive notifications. A mobile telephone number may also be provided with an account so that notifications may be sent by text message as well. Accounts are accessed using an email address or user identifier and a password.

Each user account is linked to one or more organisations. Each organisation may be associated with one or more user accounts. Each user account is limited by privileges within an organisation. Privileges are allocated to a user account by the organisation manager to allow the account to view processes and their outputs, and to execute processes. Organisations are used to separately hold processes and their data which have been executed on behalf of the organisation by account holders. Files and processes created for an organisation are not available to any other organisation.

Each organisation is associated with pre-purchased credits, which are used to pay for process executions and for the storage of files. Credits are purchased separately to the SDK. If an organisation does not have sufficient credits to execute a process, then execution will be prevented. All uploaded files and outputs from process executions incur storage costs, which are accounted for daily. If the organisation has insufficient credits to store files, then upload will be prevented. If an organisation's credits run out and it has stored files, these may be deleted to prevent further costs from being incurred without payment.

2.2 Workflow

A typical workflow for the Fusion Platform® using the SDK is as follows:

1. Login to the Fusion Platform® and select an organisation to work with

All access to the Fusion Platform® is via a registered user account. Once logged in, an organisation must be selected for which processes are to be executed. The list of organisations available to a user account can be obtained from the user information.

2. Upload a geospatial vector file which describes a region of interest

The Fusion Platform® predominately works with geospatial data, and consequently the algorithms available typically require a region of interest to be defined to produce outputs which correspond to this region. Regions of interest are defined using geospatial vector files which describe one or more polygons. The first step in executing a process is therefore to upload a vector file defining the region of interest. For example, this could be a GeoJSON file, or an ESRI Shapefile. Vector files can be created using tools such as QGIS [1], or geojson.io [2].

3. Select the service that is to be executed

The Fusion Platform® provides a variety of services which can be selected for execution. This step therefore involves obtaining the list of services available to find the one which is needed.

4. For the selected service, configure a process to run for the region of interest

Once selected, a process can be created from the service. The inputs and options for the process must be defined so that the process can be executed. Once created, and prior to execution, the process is associated with a price giving the number of credits needed to execute it. This price will be charged per execution of the process.

5. Execute the configured process

Once the process has been correctly configured to obtain a price per execution, the process can be executed. Execution can be once or scheduled to run at one or more times in the future. Each subsequent execution will incur the associated price being charged.

6. Monitor the execution of the process

Once scheduled for execution (once or many times), each execution may be monitored to determine its progress and to find out when it is finished.

For large regions of interest, the Fusion Platform® will automatically cut up the region into smaller slices to make processing more efficient. Each slice will be executed in parallel so that the outputs are obtained more quickly. When a region is cut up, multiple executions will be created automatically for a process (each with a proportion of the overall price), and all of them placed into a single group. Monitoring may therefore take place of single or grouped executions.

When the execution (or group of executions) has run to completion, an optional notification will be sent via email or text message to the user account which executed the process to notify that the execution or group of executions has finished.

It is also possible to stop all executions of a process. This will stop all ongoing and future executions. Any partial executions will be subject to the associated portion of the charges incurred.

7. Download the outputs from the executed process

When an execution has completed successfully, the outputs can be downloaded. All the inputs and outputs to an execution, and its individual steps, are available for download. The individual steps will include, for example, the ingestion of data, so that all the data used by the process is available for download.

8. Delete completed process

When a process and its outputs are no longer needed, the process can be deleted. Executions will be automatically deleted after a delay configured when the process is created. Deleting a process will delete all its executions and their outputs (and therefore stop any further storage charges being incurred). Files which have been uploaded explicitly by the user are not automatically deleted, only executions and their generated data.

9. Delete the region of interest file

When a region of interest file is no longer required, and it is not used by any process, then it can be explicitly deleted by the user (and therefore stop any further storage charges being incurred for it).

2.3

Installation

The SDK has been built for Python 3, and is tested against Python 3.7, 3.8, 3.9 and 3.10. To install the SDK into a suitable Python environment containing “pip”, execute the following:

```
pip install fusion-platform-python-sdk
```

Command

This will install the SDK and all its dependencies.

To update an existing installation to the latest version, execute the following:

```
pip install fusion-platform-python-sdk --upgrade
```

Command

2.4

Usage

Within a Python file, the SDK can then be used by importing it. The following shows how the version of the SDK can be displayed.

```
import fusion_platform

# Print the current version of the Fusion Platform(r) SDK.
print(fusion_platform.__version__)
print(fusion_platform.__version_date__)
```

Python

3 Quick Example

The following shows a simple example of how to use the SDK to create a process, execute it and download the resulting data. This example assumes that a suitable user account has been created with an organisation which has sufficient credits to run the process, and that the SDK has been setup as described in section 2.3. To use a different user account, replace the email address 'me@org.com' and password 'MyPassword123!'.

The example uses a pre-defined region of interest for Glasgow, and then obtains spectral indices for the region. A range of spectral indices are built using ESA's Sentinel-2 multi-spectral data from the most recent date available for Glasgow.

Python

```
import os

import fusion_platform

# Login with email address and password.
user = fusion_platform.login(email='me@org.com',
                             password='MyPassword123!')

# Select the organisation which will own the file. This is
# the first organisation the user belongs to in this example,
# and which is therefore assumed to have sufficient credits.
organisation = next(user.organisations)

# Create a data item for the Glasgow region of interest:
glasgow = organisation.create_data(
    name='Glasgow',
    file_type=fusion_platform.FILE_TYPE_GEOJSON,
    files=[fusion_platform.EXAMPLE_GLASGOW_FILE],
    wait=True)

# Find the spectral indices service using its name.
service, _ = organisation.find_services(name='Spectral Indices')

# Create a template process from the service.
process = organisation.new_process(name='Example',
                                   service=service)

# Configure the process to use Glasgow as the region of interest.
process.update(input_number=1, data=glasgow)

# The process defaults to run once using the most recent data.
# Create the process, which will validate its options and inputs.
process.create()
```



```
# Before execution, review the price in credits.
print(f"Price: {process.price}")

# Now execute the process and wait for it to complete.
process.execute(wait=True)

# Get the corresponding execution of the process. This is assumed
# to be the most recently started execution.
execution = next(process.executions)

# Now download all the outputs.
for component in execution.components:
    print(f"Downloading {component.name}")
    component_dir = component.name

    for output in component.outputs:
        dir = os.path.join(component_dir, output.name)

        for file in output.files:
            file.download(path=os.path.join(dir, file.file_name))

# Now tidy everything up by deleting the process.
process.delete()

# After a delay to delete the process and its executions,
# delete the uploaded file. This will fail until the executions
# have been deleted.
glasgow.delete()
```

4 Authentication

To login to the Fusion Platform® using the SDK, the user account credentials must be obtained first. These can be obtained, for example, from Digital Content Analysis Technology Ltd via support@d-cat.co.uk. The user account credentials require an email address for notifications. An optional user identifier may be supplied along with an automatically generated password. This password should be changed after first login.

To login with the SDK, use the following:

```
import fusion_platform

# Login with email address and password.
user = fusion_platform.login(email='me@org.com',
                             password='MyPassword123!')
```

Python

If you have a user identifier, you can also use:

```
# Login with user identifier and password.
user = fusion_platform.login(
    id='ff1fbd91-bd64-4ad7-ac80-d0b3f22f9d19',
    password='MyPassword123!')
```

Python

Here, the registered email address is 'me@org.com', user identifier is 'ff1fbd91-bd64-4ad7-ac80-d0b3f22f9d19' and password is 'MyPassword123!'.

If login is successful, a user object will be returned. If unsuccessful, an exception will be raised giving details of the error encountered.

The created user object provides all access to the Fusion Platform® through an authenticated session with an appropriate configuration.

Important

When using multiple threads, separate user objects should be created within each thread to ensure that no session data is changed asynchronously across threads.

When logging in, the following configuration parameters may be set in addition to the user account credentials:

`api_url` The URL of the Fusion Platform®. This will default to `https://api.thefusionplatform.com`, but may be modified as required for specific uses.

For example:

```
# Login with email address, password and custom API URL.
user = fusion_platform.login(
    email='me@org.com',
    password='MyPassword123!',
    api_url='https://custom.api.thefusionplatform.com')
```

Python

The user object represents all information about the user, as well as the active Fusion Platform® session. It can be used to modify key user details, including the user's given name, family name, telephone number and notification preferences. For example:

Python

```
# Modify the user's name:
user.update(given_name='Joe', family_name='Bloggs')
```

If the update fails for any reason, such as invalid parameter values, then an exception will be thrown. Section 0 includes full details of the content of the user information and what can be modified.

Changing the password uses a separate method since the old password must be provided:

Python

```
# Modify the password.
user.change_password(old='MyOldPassword123!',
                    new='MyNewPassword123!')
```

Note that the password must comply with the password policy enforced at the time. Currently, this requires a minimum length of 8 characters with at least one number, one lowercase letter, one uppercase letter and one of the special characters “. [] {} () ? - ! @ # % & , > < : ; | _ ~”.

The user object can also be used to provide an iterator through the organisations to which the user account belongs:

Python

```
# Get the organisation information for all the user's
# organisations.
for organisation in user.organisations:
    print(organisation.name)
```

An organisation object is used to access files and processes owned by the organisation. It can also be used to modify key organisation details if the user account has the appropriate privilege. For example:

Python

```
# Modify the organisation's name:
organisation.update(name='My Organisation')
```

If the update fails for any reason, such as invalid parameter values or insufficient privileges, then an exception will be thrown. Section 0 includes full details of the content of the organisation information and what can be modified by a suitably privileged user.

5 Uploading Files

Most algorithms available in the Fusion Platform® analyse and output data based upon a given region of interest. Regions of interest are defined using geospatial vector files which describe one or more polygons, which can be created using tools such as QGIS [1], or geojson.io [2]. Supported geospatial vector file formats include GeoJSON, ESRI Shapefile, KML and KMZ. All regions of interest are defined within a single file. For example, a GeoJSON file is a text file containing JSON which defines the region.

Important

An ESRI Shapefile is different in that the region is defined by a series of files with the same name, but different extensions. To be uploaded as a region of interest, all the files which go to make up the ESRI Shapefile must be zipped into a single archive file with no directory structure included. This single file can then be uploaded as an ESRI Shapefile.

Other algorithms may work with other types of input, such as geospatial raster or CSV files. The type of file required is detailed with the service information.

To be able to provide a region of interest file to an algorithm, the file must first be uploaded to the platform. A data object represents one or more files grouped into a single object, and data objects are used as input to, and output from, processes. For example, a region of interest defined in a GeoJSON file 'region.geojson' found in the local directory is uploaded to a data object as follows:

Python

```
# Select the organisation which will own the file. This is
# the first organisation the user belongs to in this example:
organisation = next(user.organisations)

# Now create a data object and upload the file to it.
# This will return as soon as the data object is created, and
# will not wait for the files to be uploaded or analysed.

data = organisation.create_data(name='My Region',
                                file_type=fusion_platform.FILE_TYPE_GEOJSON,
                                files=['region.geojson'])
```

This will create the associated data object for the organisation, give it a name and then upload the GeoJSON file to the Fusion Platform®. Multiple files of the same type can be uploaded by extending the list of paths in the array. If the create fails for any reason, such as invalid parameter values, insufficient privileges or not enough credits, then an exception will be thrown. The default usage of this method is not to wait for the upload and analysis of the files to complete, but instead to return once the data object has been created successfully.

Important

The amount of time it takes to upload files will depend upon the size of the files and the bandwidth of your network connection. Once uploaded, the Fusion Platform® will perform some analysis on the files before the files can be used. The default version of `create_data` will not wait for all these steps to complete, but these steps must be complete before the data object can be used within an execution.

To find out if the upload and analysis have completed, use the following:

Python

```
import time

# Loop waiting for the create to complete.
while not data.create_complete():
    time.sleep(10) # Sleep for 10 seconds.

# The create is complete without exception.
print('Create successfully completed')
```

If the upload and analysis have not yet completed, then this method will return 'False'. If it has completed successfully, this will return 'True'. If, however, the upload and analysis have completed but failed, then this method will raise an exception containing the details of why it failed.

To have the method wait for all uploads and the analysis to complete, then the 'wait' flag can be turned on:

Python

```
# Now create a data item and upload the file to it.
# This will wait for the upload and analysis to complete.
data = organisation.create_data(name='My Region',
                                file_type=fusion_platform.FILE_TYPE_GEOJSON,
                                files=['region.geojson'],
                                wait=True)
```

Important

If the upload and analysis of the data object's files has not completed, then the data object will not be available for use within a process. The upload will take place in a series of background threads, and hence if the program is closed, the upload will terminate prematurely causing the uploaded files to be in an invalid state. The data object must then be deleted, and the upload started again.

An iterator through the data objects which have been uploaded to the platform for the organisation can be obtained using:

Python

```
# Get an iterator of the data items belonging to the
# organisation:
for data in organisation.data:
    print(data.name)
```

Data items can also be searched for by name:

Python

```
# Search for data items by name:
data, data_items = organisation.find_data(name='My Region')
```

Searching by name looks for those services which have the corresponding name. Names are not unique, hence the 'data' return is the first item that has been found (or 'None' if no items found) and the 'data_items' return is an iterator through all items found.

A specific data item can also be obtained from its data identifier, which is listed as part of its information.

Python

```
# Get the data item by its identifier:
data, _ = organisation.find_data(
    id='e1fca0a4-e14d-4a9a-8c76-82a6a6831860')
```

A data object can be used to modify the name. For example:

Python

```
# Modify the name of the first uploaded data item:
data = next(organisation.data)
data.update(name='Other Region')
```

If the update fails for any reason, such as invalid parameter values or insufficient privileges, then an exception will be thrown. Section 0 includes full details of the content of the data object and what can be modified.

A data object can be used to iterate through and download one or more of the associated files:

Python

```
# Download all files from the data object to the local
# directory. Each download will return immediately once started.
for file in data.files:
    print(f"Downloading {file.name} ({file.size} bytes)")
    file.download(path=file.name)
```

Section 0 includes full details of the content of the file information.

The default usage of the 'download' method is not to wait for the for the download to complete, but instead to return immediately.

Important

The amount of time it takes to download a file will depend upon the size of the files and the bandwidth of your network connection. The download will take place in a background thread, and hence if the program is closed, the download will terminate prematurely causing the downloaded file to be in an invalid state.

To find out the progress of a download, and if the download has completed, use the following:

Python

```
import time

# Loop waiting for the download to complete.
while not file.download_complete():
    url, destination, size = file.download_progress()
    print(f"Downloaded {size} of {file.size} bytes")

    time.sleep(10) # Sleep for 10 seconds.

# The download is complete without exception.
print('Download successfully completed')
```

To have the method wait for the download to complete, turn the 'wait' flag on:

Python

```
# Download all files from the data object to the local
# directory. Each download will only return when complete.
for file in data.files:
    print(f"Downloading {file.name} ({file.size} bytes)")
    file.download(path=file.name, wait=True)
```

When the data item is not being used by any process, it may be deleted:

Python

```
# Delete the data item:
data.delete()
```

This will raise an exception if the data item is in use by a process, or because of insufficient privileges, otherwise the item will be deleted. Deleting a data item ensures that no further storage costs are incurred for its files. Storage charges are incurred daily.

6 Browsing Services

The purpose of the Fusion Platform® is to provide a convenient way of efficiently executing a range of algorithms. Each algorithm is bundled into a service which defines what data is required for the algorithm to execute so that when the service is run, it will obtain all the required data, then run the algorithm to produce the desired outputs.

The following provides an iterator through the services that are available to an organisation:

Python

```
# Select the organisation to browse services. This is
# the first organisation the user belongs to in this example:
organisation = next(user.organisations)

# Now iterate through the available services:
for service in organisation.services:
    print(service.name)
```

A service object contains all the details of a service. This includes its name and documentation. Services can also be searched for by name or by keyword:

Python

```
# Search for services by name:
service, services = organisation.find_services(
    name='Sentinel-2 (Multispectral)')

# Search for services by keyword:
service, services = organisation.find_services(keyword='cloud')
```

Searching by name looks for those services which have the specified text in their name, whereas a keyword search will search service keywords to find any which contain the corresponding keyword. All searches are case sensitive. Names are not unique, hence the 'service' return is the first item that has been found (or 'None' if no items found) and the 'services' return is an iterator through all items found.

A specific service can also be obtained from its service identifier or SSD ID. Both are listed as part of the service information.

Python

```
# Both these get the same Sentinel-2 service:
service, _ = organisation.find_services(
    id='5e98fdd7-77a3-451f-ba49-d67002fc420d')
service, _ = organisation.find_services(
    ssd_id='819f9e0c-a94e-4477-81f6-467a0f876e77')
```

Section 0 includes full details of the content of the service information.

Important

All services in the Fusion Platform® are uniquely identified by an SSD ID. At any time, there will always be only one service available with the unique SSD ID. When a new version of a service is released, it will have the same SSD ID, but it will have a different version number and individual service identifier. The service identifier therefore uniquely identifies a specific version of a service, whereas an SSD ID is used to find the most recent version of a service, and should therefore be the preferred way of finding a service through an identifier.

7 Creating Processes

Once a service has been selected, a process can be created from it. To obtain a template process from a service, use the following:

Python

```
# Get a template process from a service object:
process = organisation.new_process(name='Example',
                                   service=service)
```

Here, to create the template process, a service object is required. If the request fails for any reason, such as invalid parameter values or insufficient privileges, then an exception will be thrown. See section 6 for a description of how to get a service object.

Important

The returned process object is a template which has not been configured for execution. Because it has not been configured, this template object has not been persisted to the Fusion Platform®. Only a successfully configured process object can be persisted.

The process object contains three important parts that must be completed before the process object can be saved to the Fusion Platform®. First, all mandatory options must be completed. Second, all input data items need to be specified. Third the execution schedule needs to be defined.

Section 0 includes full details of the content of the process information. Options can be completed as per the following example:

Python

```
from datetime import datetime, timezone

# Set the mandatory options which need to be completed:
for option in process.options:
    if option.required and (value is None):
        # This option is mandatory. Set the value by its type.
        if option.data_type == fusion_platform.DATA_TYPE_NUMERIC:
            process.update(option=option, value=1)
        elif option.data_type == fusion_platform.DATA_TYPE_CURRENCY:
            process.update(option=option, value=1)
        elif option.data_type == fusion_platform.DATA_TYPE_BOOLEAN:
            process.update(option=option, value=True)
        elif option.data_type == fusion_platform.DATA_TYPE_DATETIME:
            process.update(option=option, value= datetime.now(timezone.utc))
        elif option.data_type == fusion_platform.DATA_TYPE_STRING:
            process.update(option=option, value='Hello world!')
        elif option.data_type == fusion_platform.DATA_TYPE_CONSTRAINED:
            process.update(option=option, value=option.constrained_values[0])
```

This example loops through all the options, finds the mandatory options which have not got a default value, and completes each with an example value of the correct data type. Alternatively, individual options can be set using their name:

Python

```
from datetime import datetime, timedelta, timezone

# Set an option using its name.
process.update(option_name='latest_date', value=False)

start_date = datetime.now(timezone.utc) - timedelta(days=25)
process.update(option_name='start_date', value=start_date)
```

To provide inputs, there must be an existing data model which has been uploaded and completed its analysis, as shown in section 5.

Python

```
# Find the region of interest and a raster file, both of
# which have been uploaded previously.
region, _ = organisation.find_data(name='My Region')
raster, _ = organisation.find_data(name='My Raster')

# Specify the inputs to the process.
for input in process.inputs:
    if input.file_type == fusion_platform.FILE_TYPE_GEOJSON:
        process.update(input=input, data=region)
    elif input.file_type == fusion_platform.FILE_TYPE_GEOTIFF:
        process.update(input=input, data=raster)
```

Alternatively, inputs can be set using their input number, with the number for the first input to the process '1'.

Python

```
# Set the input by its number.
process.update(input_number=1, data=region)
```

Finally, the execution schedule can be defined. This is either to run the process once, or to run it using a particular schedule:

Python

```
from datetime import datetime, timedelta, timezone

# Run the process once...
process.update(run_type=fusion_platform.RUN_TYPE_RUN_ONCE)

# ...or using a schedule, which will run the job once, tomorrow.
then = datetime.now(timezone.utc) + timedelta(days=1)
process.update(run_type=fusion_platform.RUN_TYPE_RUN_SCHEDULE,
               repeat_count=1, repeat_start=then)
```

Full details on options, inputs and scheduling can be found with the process details in section 0.

When the process has been correctly configured, it can be saved to the Fusion Platform®. During this creation, the process will be validated and its corresponding price per execution calculated. If the process fails validation, then an exception will be thrown, and the process will not be saved.

Python

```
# Attempt to create the process which will first validate it.
# This will throw an exception if the validation fails.
process.create()

# Once created, the price will be available to review.
print(process.price)
```

Important

Prior to execution, the name and schedule of a process may be changed using the 'update' method. However, the inputs and options may only be configured prior to creating the process.

When the process is no longer needed, it may be deleted:

Python

```
# Delete the process:
process.delete()
```

This will raise an exception in case of insufficient privileges, otherwise the process will be deleted. This will first stop any current executions and prevent any further executions taking place. Then this will delete any existing executions and their associated outputs, before deleting the process itself. This deletion cannot be undone.

8 Executing and Monitoring Processes

An iterator through the processes which have been validated and saved for an organisation can be obtained as follows:

Python

```
# Select the organisation to find processes. This is
# the first organisation the user belongs to in this example:
organisation = next(user.organisations)

# Now iterate through the processes:
for process in organisation.processes:
    print(f"{process.name}: {process.process_status}")
```

Each entry is a process object, which contains all the details of a created process. Processes can also be searched for by name:

Python

```
# Search for processes by name:
process, processes = organisation.find_processes(name='Process')
```

Searching by name looks for those processes which have the corresponding name. All searches are case sensitive. Names are not unique, hence the 'process' return is the first item that has been found (or 'None' if no items found) and the 'processes' return is an iterator through all items found.

A specific process can also be obtained from its process identifier, which is listed as part of its process information.

Python

```
# Get the process by its identifier:
process, _ = organisation.find_processes(
    id='6a4417de-e885-4749-9f75-f7eeb2aac7ab')
```

Section 0 includes full details of the content of the process information.

A process may be executed using the following:

Python

```
# Execute the configured process:
process.execute()
```

This will execute a process according to its options, inputs and schedule. Every execution will incur the price listed with the process. This method will return immediately and will not wait for the execution to complete.

To wait for the execution to start after calling 'execute', use the following:

Python

```
# Wait for the next execution to start.
process.wait_for_next_execution()
```

This will block until the next execution has started.

Important

If the process is configured with a schedule, this method will wait until the scheduled time, and hence this method may block for a considerable time.

For every execution of a process, a corresponding process execution object is created, and these can be obtained as follows:

Python

```
# Get all the executions for the process:
for execution in process.executions:
    print(f"{execution.id}: {execution.progress}")
```

The progress of an execution can be determined from its 'progress' attribute, which is its percentage progress through execution. When the 'progress' is 100%, then the execution is complete and the 'success' flag will state whether the execution was successful ('True') or failed ('False'). When an execution completes, an optional notification will be sent by email and/or text message to the user account that executed the process. Whether such notifications are sent can be configured for the user account (see section 4 for how to modify account information).

For convenience, the following can be used to determine if the execution has finished, raising an exception if the execution was unsuccessful:

Python

```
import time

# Loop waiting for the execution to complete.
while not execution.check_complete():
    time.sleep(10) # Sleep for 10 seconds.

# The execution has completed without exception.
print('Execution successfully completed')
```

To have the method wait for the execution to complete, turn the 'wait' flag on:

Python

```
# Wait for the execution to complete.
execution.check_complete(wait=True)

# The execution has completed without exception.
print('Execution successfully completed')
```

To execute and wait, the 'wait' flag can be turned on:

Python

```
# Execute the configured process and wait for completion:
process.execute(wait=True)
```

Important

The amount of time it takes to execute the process depends upon the selected service. When the 'wait' flag is turned on, the method will wait for the next scheduled execution to start, and then block until all current executions of the process have completed. If the process is configured to repeat the execution more than once, only the first execution will be considered.

For large regions of interest, the Fusion Platform® will automatically cut up the region into smaller slices to make processing more efficient, and each slice will be allocated to a separate execution within a group. If an execution is part of a group, then its 'group_id' value will be set. Just the executions for a group can be obtained as follows:

Python

```
# Get all the executions for a group:
execution, executions = process.find_executions(
    group_id='1cf57d5c-f1c9-4cac-b6c7-076cc9c3ec89')
```

The 'execution' return is the first execution that has been found (or 'None' if no items found) and the 'executions' return is an iterator through all items found. When an execution is part of a group, only one completion notification will be sent, instead of one per execution in the group.

Every process is configured with an output storage period in days which determines how long each execution is kept before it is automatically deleted. Prior to deletion, an optional notification may be sent via email or text message to the user account which executed the process to warn that the execution will be deleted shortly.

If the automatic deletion should be delayed further, then this delete expiry can be modified for an execution as follows:

Python

```
# Change the delete expiry for an execution to delay it
# from being deleted for 2 days from now.

delete_expiry = datetime.now(timezone.utc) + timedelta(days=2)
execution.update(delete_expiry=delete_expiry)
```

If the update fails for any reason, such as invalid parameter values or insufficient privileges, then an exception will be thrown. Extending the expiry will mean that additional storage costs are incurred for the extra time that the execution's outputs are stored. All storage charges are collated daily. Section 0 includes full details of the content of the execution object and what can be modified.

Process executions can be stopped at any time by using the following:

Python

```
# Stop all executions of a process, aborting those running:

process.stop()
```

This will first stop the process from running any more executions as per its schedule. This will then abort any executions which are in progress. Partial charges will still be incurred for any incomplete execution and the corresponding outputs which have been created.

9 Obtaining Execution Inputs and Outputs

Once an execution has completed, then its inputs and outputs can be downloaded. Each execution is composed of the different operations that have been executed to make up the full processing chain of the service. For example, this will involve ingesting data and pre-processing it before the associated algorithm is run. Each of these components therefore corresponds to a separate, individual service executed as part of the chain. Each of these components can be obtained from an execution as follows:

Python

```
# Select the execution for a process. This is the last
# execution which was started for the process in this example:
execution = next(process.executions)

# Now iterate through the components of the execution:
for component in execution.components:
    print(f"{component.name}: {component.runtime}")
```

If an execution is still executing, then it and its components can still be accessed using the above methods. However, the associated data will be incomplete. Section 0 includes full details of the content of the execution component information.

Each component may have inputs and outputs associated with it. Typically in a processing chain, the first component in the chain will take as input the region of interest. Its outputs will then be fed as inputs to the second component in the chain, and so on. Whereas a typical first input is a region of interest defined as a geospatial vector file (see section 5), outputs from executions may be any supported file type. This includes geospatial vector files, such as GeoJSON, ESRI Shapefile, KML and KMZ and geospatial raster files, including DEM, GeoTIFF, JPEG2000. Non-geospatial images, such as PNG and JPEG files, as well as other types of file may also be output, including CSV files. Types of file unknown to the Fusion Platform® will be listed as 'Other'.

All the inputs and outputs are data objects which are associated with one or more files each. (See section 5 for how to work with data objects.)

Each of the files associated with the input data objects can be downloaded as follows:

Python

```
# Select the first input for a component in this example:
input = next(component.inputs)

# Download all files from the input to the local directory:
for file in input.files:
    print(f"Downloading {file.name} ({file.size} bytes)")
    file.download(path=file.name)
```

Similarly, the files from an output can also be downloaded:

Python

```
# Select the first output for a component in this example:
output = next(component.output)

# Download all files from the output to the local directory:
for file in output.files:
    print(f"Downloading {file.name} ({file.size} bytes)")
    file.download(path=file.name)
```

Executions will be automatically deleted according to the output storage period defined for the associated process (see section 8). However, an execution can be deleted before this expiry as follows:

Python

```
# Delete the execution:
execution.delete()
```

This will raise an exception in case of insufficient privileges, otherwise the execution and its associated outputs will be deleted. This deletion cannot be undone.

10 Detailed Example

The following shows a detailed example of how to use the SDK to create a process, execute it and download the resulting data. This example assumes that a suitable user account has been created with an organisation which has sufficient credits to run the process, and that the SDK has been setup as described in section 2.3. To use a different user account, replace the email address 'me@org.com' and password 'MyPassword123!'.

The example uses a pre-defined region of interest for the Lake District, and then performs cloud classification over the region for a specific date. Since the region is relatively large, its execution is split into four areas. The example shows one way in which these split results can be combined into a single GeoTIFF, as well as how to work with the metadata generated by the Fusion Platform® for each file.

Python

```
from datetime import datetime, timezone
import os
import rasterio
from rasterio.merge import merge
from rasterio.plot import show
import matplotlib.pyplot as pyplot

import fusion_platform

# Login with email address and password.
user = fusion_platform.login(email='me@org.com',
                             password='MyPassword123!')

# Select the organisation which will own the file. This is
# the first organisation the user belongs to in this example,
# and which is therefore assumed to have sufficient credits.
organisation = next(user.organisations)

# Create a data item for the Lake District region of interest:
lake_district = organisation.create_data(
    name='Lake District National Park',
    file_type=fusion_platform.FILE_TYPE_GEOJSON,
    files=[fusion_platform.EXAMPLE_LAKE_DISTRICT_FILE],
    wait=True)

# Find the cloud classification service using its name.
service, _ = organisation.find_services(name='Cloud
                                         Classification')

# Create a template process from the service.
process = organisation.new_process(name='Example',
                                   service=service)
```

```

# Configure the process to use Lake District National Park
# as the region of interest.
process.update(input_number=1, data=lake_district)

# Configure the service to work on historic data.
process.update(option_name='latest_date', value=False)
process.update(option_name='start_date',
                value=datetime(2021, 3, 27, tzinfo=timezone.utc))
process.update(option_name='end_date', value=None)

# Create the process, review its price, then execute it.
process.create()
print(f"Price: {process.price}")
process.execute(wait=True)

# Because the Lake District region of interest is relatively
# large, the process will be executed in a group, so find all of
# the executions in the group. These are assumed to be the most
# recently started group obtained from the last execution.
execution = next(process.executions)
_, executions = process.find_executions(
    group_id=execution.group_id)

# Download the cloud classification output for each execution.
# We also keep the metadata for each raster file.
files = []
selectors = []

for execution in executions:
    for index, component in enumerate(execution.components):
        # We are only interested in the first file for each cloud
        # classification output.
        if component.name == 'Cloud Classification':
            print(f"Downloading {component.id}")
            output = next(component.outputs)
            file = next(output.files)
            files.append(file)
            file.download(path=f"{index}_{file.file_name}")

        # The file has a single band for the classification.
        selectors.append(file.selectors[0])

```

```

# Wait for all downloads to complete.
for file in files:
    file.download_complete(wait=True)

# Create a single cloud classification image as a mosaic.
rasters = []

for file in files:
    rasters.append(rasterio.open(file))

mosaic, transform = merge(rasters)
mosaic_meta = rasters[0].meta.copy()
mosaic_meta.update({'height': mosaic.shape[1],
                    'width': mosaic.shape[2],
                    'transform': transform})

with rasterio.open('mosaic.tif', 'w', ** mosaic_meta) as merged:
    merged.write(mosaic)

# Display the mosaic image, and the histogram.
with rasterio.open('mosaic.tif', 'r') as source:
    pyplot.imshow(source.read(1, masked=True), cmap='coolwarm')
    pyplot.show()

histogram = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

for metadata in selectors:
    histogram = [first + second for first, second in
                 zip(histogram, metadata['histogram'])]

# There are only 4 classes, while the histogram has 10 bins.
x = ['Ground', 'Shadow', 'Thick Cloud', 'Thin Cloud']
y = [histogram[0], histogram[3], histogram[6], histogram[9]]

pyplot.bar(x, y)
pyplot.show()

# Tidy everything up by deleting the process and the input file.
process.delete()
lake_district.delete()

```

11 SDK Object Attributes and Methods

The SDK consists of package functions which control the overall behaviour of the SDK, and model objects, which are created in response to requests to the Fusion Platform®. Key model objects include:

User	The current logged in user associated with the active session with the Fusion Platform®.
Organisation	An organisation that the user belongs to. Organisations own data and processes.
Credit	The credits assigned to an organisation.
Data	A data item which holds one or more files of a specific type.
File	A data object's file and its associated metadata.
Service	An available service which may be chosen to be configured as a process.
Process	Describes the specific inputs and options for a service which may be executed to produce output data objects.
Execution	A specific execution of a process.
Component	Each service is composed of individual components used to complete the desired processing. For example, the ingestion of data followed by its classification. Each execution is therefore made up of one or more execution components representing these stages of processing.
Log	The log output from the execution of a component.

Each model object has three elements:

Attributes	These represent the information associated with each object, such as a 'name'. Attributes may also be used to obtain an iterator through child objects, such as 'executions' for a process. All attributes are read-only.
Common Methods	Common methods include 'update', which, where allowed, permit one or more of the object's attributes to be updated. For example, 'process.update(name="My process")'. Other common methods include 'find_<object_type>' and 'delete', where relevant and permitted.
Specific Methods	Specific methods are only available to certain object types. For example, 'organisation.new_process(service=service)', which allows a template process to be initialised for an organisation.

11.1 Attributes and Methods

Detailed information about the available attributes and methods for the package and its classes used to form model objects can be obtained from:

https://www.d-cat.co.uk/public/fusion_platform_python_sdk/

Similar information can also be obtained in Python using the following:

Python

```
import fusion_platform

# Get help on the package methods:
help(fusion platform)
```

Help on individual classes can be obtained in a similar way:

Python

```
from fusion_platform.models.user import User

# Get help on the user class:
help(User)
```

12 Glossary of Terms

Application Programming Interface (API)	Typically refers to a software interface (rather than a user interface) used to access services or resources, performing actions or storing or retrieving data.
Comma Separated Values (CSV)	A text file format used to provide row-based data separated by commas.
Digital Elevation Model (DEM)	A type of three-dimensional model file for geolocated elevations.
Earth Observation (EO)	Gathering of physical characteristics of the Earth, typically via remote sensing.
ESRI Shapefile	A type of vector data file which allows attributes to be assigned to locations using points, lines or polygons.
GeoJSON	JSON data used to encode geospatial data.
Geospatial raster file	A file which contains raster data which is geospatially located. GeoTIFF and JPEG2000 are both types of geospatial vector file.
Geospatial vector file	A file which is used to describe geographic areas or points of interest. GeoJSON and ESRI Shapefiles are both examples of geospatial vector files.
GeoTIFF	A type of image file which allows the embedding of geospatial metadata alongside image data (in TIFF format).
Hypertext Transfer Protocol (HTTP)	Communication protocol used to exchange data with a web server. In this document, HTTP is synonymous with the corresponding HTTP Secure version (HTTPS).
Joint Photographic Experts Group (JPEG)	A type of image file format which uses lossy data compression.
JPEG2000	A type of image file which allows the embedding of geospatial metadata alongside image data (in JPEG format).
JavaScript Object Notation (JSON)	Text-based format used to label data for data interchange.
Keyhole Mark-up Language (KML)	A text-based mark-up language used to encode geospatial data, with the zipped version being KMZ.
Keyhole Mark-up Language Zipped (KMZ)	A KML file and associated resources which has been zipped into a single file to reduce the file size.
OpenAPI 3.0 Specification	The standardised specification of a RESTful API interface which can be used to document or access the API with suitable software which can read the specification.

Portable Network Graphics (PNG)	A type of image file format which uses lossless data compression.
Process	Refers to process which has been created from a selected service, and which defines the required inputs and options for the process to execute.
Remote Sensing	Refers to the use of satellite (or aircraft or drone) technologies to remotely sense the Earth. A variety of sensing technologies can be used, including optical, hyperspectral or radar sensors.
Service	Services within the Fusion Platform® can be selected to execute a particular algorithm given appropriate inputs and options.
Software Development Kit (SDK)	Term used to define software used by developers to perform a specific function. In this document, the SDK is used to interact with the Fusion Platform®.
Standardised Service Description (SSD)	Within the Fusion Platform® this defines the input, algorithm options and output requirements of a service which may be selected and configured as a process to be run. An SSD ID uniquely identifies a service which can be selected.
Uniform Resource Locator (URL)	A reference to a resource held on the web. A URL consists of a protocol (for example, "https://") a resource name ("thefusionplatform.com"), a path ("/platform") and parameters ("?id=1&name=User").
Zip Archive	A single file which contains other files which have been compressed.

13 References

- [1] QGIS, “Welcome to the QGIS project!,” 2022. [Online]. Available: <https://www.qgis.org>. [Accessed February 2022].
- [2] Mapbox, “geojson.io,” 2022. [Online]. Available: <https://geojson.io>. [Accessed February 2022].